

SafeLog AI — Architecture Report

Author: Szymon Iwacz · **Stack:** Rails 8.1 monolith, SQLite, Devise, server-rendered ERB · **Live demo:** <https://safelog-ai.fly.dev/>

1. System overview

SafeLog AI helps engineers debug production issues from **multiple pasted log sources** without leaking secrets. The core insight: **deterministic redaction runs in memory before any persistence or AI reasoning**. Raw pasted content and raw→placeholder mappings exist only for the duration of the request.

```
User → DebuggingCasesController → Intake::ProcessCaseSubmission
      → Redaction::Engine → SQLite (encrypted)
Controller → Analysis::AnalyzeCase → Correlation::ExtractSignals
                                   → Ai::Client (via PromptBuilder) → AiReport
Boundary: Redaction must not import Ai
```

Runtime territory (M4L2): product logic lives in `app/services/{intake,redaction,correlation,analysis,ai}/`; HTTP is thin; security oracles in `spec/requests/*_security*`. High git activity in `context/changes/` is **documentation workflow**, not deployable runtime — a common catalog illusion in this repo.

Structure (M4L2): Ruby services form a **DAG with no cycles**. Critical boundary: `redaction ⊥ ai` — verified by constant scan and dependency analysis (`artifact-2-structure.md`). E2E tests are black-box (`depcruise:validate` — 0 violations).

2. Intake flow and refactor direction (M4L3–M4L4)

Primary flow researched: `POST /debugging_cases` → `Intake::CaseSubmission` (validation) → `Intake::ProcessCaseSubmission` (transaction, redaction, persist) → show page with sanitized evidence and redaction summary.

Findings (ast-grep verified):

- Single runtime path from paste to persist; no parallel intake shortcuts in production.
- `redaction_findings.create!` has one call-site in `ProcessCaseSubmission` — implicit contract between `Redaction::Engine` output and persistence.
- Security oracles prove raw substrings never land in SQLite; AI prompts contain placeholders only.

Ranked refactor opportunities (exploration only — M4L4):

Rank	ID	Opportunity	Why first
1	TD-2	Typed finding persist boundary	Lowest cost; closes implicit hash contract (partially shipped post-M4 as <code>RedactionFinding.build_from_engine_finding</code>)
2	IMPL-1	Extract persist object from <code>ProcessCaseSubmission</code>	Composes with TD-2; enables rollback specs

Rank	ID	Opportunity	Why first
3	TD-5	CRLF normalization in <code>Redaction::Engine</code>	Low blast radius; paste correctness

Hard rule during M4: **research and ranking only** — no mandatory production refactor for the Architect badge.

3. Domain model and planned hardening (M4L5)

Ubiquitous language: Debugging case, log source, redaction finding, correlation signal, hypothesis report. Subdomains: **Intake & Redaction** (core), **Correlation** (pure extraction), **Analysis & AI** (hypothesis generation behind adapter).

Primary invariant — INV-G1: No diagnostic content reaches persistence until it passes in-memory redaction; raw paste and mappings are never stored. Today enforced procedurally in `ProcessCaseSubmission` plus schema absence of `raw_*` columns and test oracles — **not** by a structural aggregate type.

Planned improvements (plan-only, not certification blockers):

`SanitizedCaseDraft` **aggregate guardian** — typed gate so AR models cannot accept unsanitized strings (`02-invariant-aggregate-refactor.md`).

`HypothesisGenerator` **port + ACL** — isolate OpenAI/provider details from domain; `Analysis::AnalyzeCase` depends on port, adapter wraps `Ai::Client` (`03-anti-corruption-layer.md`).

Both plans include phased F1–Fn roadmaps for post-MVP hygiene when product prioritizes structural debt over new features.

4. Quality, security, and team automation

Layer	Mechanism
Tests	256 RSpec + 9 Capybara system + 19 Playwright E2E (4 capture + 15 functional); fake AI in CI; SimpleCov 100% line + branch
CI	<code>bin/ci</code> parity with GHA (RuboCop, Brakeman, bundler-audit, importmap, RSpec)
Encryption	Active Record Encryption on sanitized logs, reports, correlation payloads
Champion (M5)	TypeScript PR review agent + GHA AI review workflow; <code>@szymoniwacz/ai-toolkit</code> on GitHub Packages

Security narrative for reviewers: transient raw intake → encrypted sanitized evidence only → scoped `current_user.debugging_cases.find` (404 cross-user) → `PromptBuilder` sanitized-only → hypothesis-framed output with uncertainty notes.

5. Evidence index

M4 lesson	Artifact
M4L2 Repo map	<code>context/map/repo-map.md</code>
M4L3 Flow research	<code>10x-archive/case-submission-flow-analysis/research.md</code>
M4L4 Refactor ranking	<code>context/changes/refactor-opportunities/research.md</code>
M4L5 Domain	<code>context/domain/</code>
Readiness audit	<code>context/reviews/m4-architect-readiness-review.md</code>
Screenshots	<code>screenshots/architect/</code>

Defensibility: This report synthesizes artifacts produced across M4L2–L5 with ast-grep verification and live code inspection — not a single-prompt dump. Next structural slice when prioritized: follow F1–F5 in `02-invariant-aggregate-refactor.md` (INV-G1 aggregate guardian).

PDF export: `architecture-report.pdf` — regenerate with `npm run cert:architecture-pdf`.